

DataCamp EDA in R

DataCamp EDA in R

README

- This is an example of how I work through a [DataCamp lesson](#) - no matter what the subject. It's a little more work when there are videos involved. Time to completion about 2 hours.
- I am pretty sure that you won't really understand what's going on without doing it at this level of granularity and with the original data set.
- This is also a way of keeping the DataCamp lessons interesting because otherwise I'd sometimes get very bored with the pace and/or the subject matter. Gotta keep things buzzing!
- As a classroom assignment, this file has been sanitized for use as a practice exercise. The complete file is available as a PDF in the `ds205/pdf/` directory in GitHub.

Emacs setup (optional)

Hide emphatic characters like ~, *

To **not** see the emphatic characters like ~ or * or / in the Org file text, run the following code chunk (or put the code in your `/.emacs` file): if successful, you should see "t" in the minibuffer.

```
(setq-default org-hide-emphasis-markers t)
```

This will only work for new buffers. If you don't put it in your `/.emacs` file, the command will only work for the current Emacs session.

Close and reopen this file to see an effect.

Change your theme

- In Emacs, type `M-x custom-themes`
- In the buffer that appears, select Leuven
- Select Apply and Save Setting
- This will work immediately

Exploring categorical data

Import and view the data

- [X] Download the raw data Comics data [from DataCamp](#) to your PC into a directory `./data/`
- [X]

Read the data set into an R data frame named `comics` and check the structure of the data frame.

```
comics <- read.csv(file="data/comics.csv", header=TRUE)
str(comics)
```

```
'data.frame': 23272 obs. of 11 variables:
 $ name      : chr  "Spider-Man (Peter Parker)" "Captain America (Steven Rogers)" "Wolverine (James \"Logan\" Howlett)" "Iron Man (Anthony \"Tony\" Stark)" "Thor (Thor Odinson)" "Benjamin Grimm (Earth-616)" "Reed Richards (Earth-616)" "Hulk (Robert Bruce Banner)" "Scott Summers (Earth-616)" "Jonathan Storm (Earth-616)"
 $ id        : chr  "Secret" "Public" "Public" "Public" ...
 $ align     : chr  "Good" "Good" "Neutral" "Good" ...
 $ eye       : chr  "Hazel Eyes" "Blue Eyes" "Blue Eyes" "Blue Eyes" ...
 $ hair      : chr  "Brown Hair" "White Hair" "Black Hair" "Black Hair" ...
 $ gender    : chr  "Male" "Male" "Male" "Male" ...
 $ gsm       : chr  NA NA NA NA ...
 $ alive     : chr  "Living Characters" "Living Characters" "Living Characters" "Living Characters" ...
 $ appearances : int  4043 3360 3061 2961 2258 2255 2072 2017 1955 1934 ...
 $ first_appear : chr  "Aug-62" "Mar-41" "Oct-74" "Mar-63" ...
 $ publisher  : chr  "marvel" "marvel" "marvel" "marvel" ...
```

- []

Print the first 10 rows of the dataset.

```
head(comics, 10)
```

	name	id	align	eye	hair
1	Spider-Man (Peter Parker)	Secret	Good	Hazel Eyes	Brown Hair
2	Captain America (Steven Rogers)	Public	Good	Blue Eyes	White Hair
3	Wolverine (James \"Logan\" Howlett)	Public	Neutral	Blue Eyes	Black Hair
4	Iron Man (Anthony \"Tony\" Stark)	Public	Good	Blue Eyes	Black Hair
5	Thor (Thor Odinson)	No Dual	Good	Blue Eyes	Blond Hair
6	Benjamin Grimm (Earth-616)	Public	Good	Blue Eyes	No Hair
7	Reed Richards (Earth-616)	Public	Good	Brown Eyes	Brown Hair
8	Hulk (Robert Bruce Banner)	Public	Good	Brown Eyes	Brown Hair
9	Scott Summers (Earth-616)	Public	Neutral	Brown Eyes	Brown Hair
10	Jonathan Storm (Earth-616)	Public	Good	Blue Eyes	Blond Hair

	gender	gsm	alive	appearances	first_appear	publisher
1	Male	<				
	Living Characters	4043	Aug-62	marvel		
2	Male	<				
	Living Characters	3360	Mar-41	marvel		
3	Male	<				
	Living Characters	3061	Oct-74	marvel		
4	Male	<				
	Living Characters	2961	Mar-63	marvel		
5	Male	<				
	Living Characters	2258	Nov-50	marvel		
6	Male	<				
	Living Characters	2255	Nov-61	marvel		
7	Male	<				
	Living Characters	2072	Nov-61	marvel		
8	Male	<				
	Living Characters	2017	May-62	marvel		
9	Male	<				
	Living Characters	1955	Sep-63	marvel		
10	Male	<				
	Living Characters	1934	Nov-61	marvel		

- []

What happens when you only enter the name of the data frame in the console? What is the default maximum value of rows displayed? Enter the command below to print this value.

```
getOption("max.print")
```

```
[1] 99999
```

- []

Unlike the data shown in the video, the data frame you are currently working with is not a `data.frame` but a `"tibble"`.

More importantly, the tibble contains factor variables where the comics data frame has character variables.

Re-import `comics.csv` into `comics` so that the characters in the data frame become factor variables. *Tip: check the `read.csv` help.*

```
comics <- read.csv(file="data/comics.csv",
                  header=TRUE,
                  stringsAsFactors=TRUE )
str(comics)
```

```
'data.frame': 23272 obs. of 11 variables:
 $ name      : Factor w/ 23272 levels "'Spinner (Earth-616)",...: 19830 3335 22769 9647 2
 $ id        : Factor w/ 4 levels "No Dual","Public",...: 3 2 2 2 1 2 2 2 2 2 ...
 $ align     : Factor w/ 4 levels "Bad","Good","Neutral",...: 2 2 3 2 2 2 2 2 3 2 ...
 $ eye       : Factor w/ 26 levels "Amber Eyes","Auburn Hair",...: 11 5 5 5 5 5 6 6 6 5 .
 $ hair      : Factor w/ 28 levels "Auburn Hair",...: 7 27 3 3 4 14 7 7 7 4 ...
 $ gender    : Factor w/ 3 levels "Female","Male",...: 2 2 2 2 2 2 2 2 2 2 ...
 $ gsm       : Factor w/ 6 levels "Bisexual Characters",...: NA NA NA NA NA NA NA NA NA
 $ alive     : Factor w/ 2 levels "Deceased Characters",...: 2 2 2 2 2 2 2 2 2 2 ...
 $ appearances : int 4043 3360 3061 2961 2258 2255 2072 2017 1955 1934 ...
 $ first_appear: Factor w/ 1606 levels "1935, October",...: 874 1278 1513 1296 1424 1432 14
 $ publisher  : Factor w/ 2 levels "dc","marvel": 2 2 2 2 2 2 2 2 2 2 ...
```

Working with factors - levels, NAs, contingency

- []

Let's look at factor levels.

- Store the variables `align` and `id` of the data frame in vectors of that name.
- Print the levels of the `align` and `id` columns that indicate how good and hidden a superhero is.
- Can you do this in 3 lines (instead of 4) using vectorization?

```
align <- comics$align
id <- comics$id
levels(c(align,id)) ## using vectorization
```

```
[1] "Bad"           "Good"          "Neutral"
[4] "Reformed Criminals" "No Dual"       "Public"
[7] "Secret"        "Unknown"
```

- []

In the video, you're told that `levels` has ignored NA values.

How many NA values does `align` and `id` have?

Tip: use the `is.na` function to find out

```
sum(is.na(align))
sum(is.na(id))
```

```
[1] 3413
[1] 5783
```

- []

Print the contingency table for `align` and `id`, which shows how these two categorical variables are connected.¹

```
tbl <- table(align, id)
tbl
```

	id				
align	No	Dual	Public	Secret	Unknown
Bad	474	2172	4493	7	
Good	647	2930	2475	0	
Neutral	390	965	959	2	
Reformed Criminals	0	1	1	0	

- []

How many good superheroes are there, who are also good? The answer, from the table, is 4493. Which command gets you this number, too?

```
names(table(align == "Bad" & id == "Secret")) ## find the name
table(align == "Bad" & id == "Secret")[2] ## extract element
```


```
[1] "FALSE" "TRUE"
TRUE
4493
```

Barplots

- [] Create a directory `./img/` for the plots.
- []

Make a barplot of the superhero identities (`x = comics$id`).


```
library(ggplot2)
ggplot(data = comics, aes(x = id)) +
  geom_bar() +
  xlab("Superhero identities comics$id")
```



- []

Make a *stacked* barplot that shows superhero identities ($x = \text{comics}\$id$) and superhero goodness ($\text{fill} = \text{comics}\$align$).


```
library(ggplot2)
ggplot(data = comics, aes(x = id, fill = align)) +
  geom_bar()
```



- []

Make a *side-by-side* barplot that shows superhero identities (`x = comics$id`) and superhero goodness (`fill = comics$align`).


```
library(ggplot2)
ggplot(data = comics, aes(x = id, fill = align)) +
  geom_bar(position = "dodge") +
  xlab("Superhero identities comics$id")
```



- []

Make a *stacked* barplot that shows the *proportion* (not the *count*) of superhero identities ($x = \text{comics}\$id$) and superhero goodness ($\text{fill} = \text{comics}\$align$).

```
library(ggplot2)
ggplot(data = comics, aes(x = id, fill = align)) +
  geom_bar(position = "fill") +
  xlab("Superhero identities comics$id") +
  ylab("Proportion")
```



- [] Go back over the last three plots and add text:
 - x labels (barplot bar1, stacked / side-by-side bar2, bar3)
 - y labels (proportional barchart bar4)
 - plot titles

Footnotes:

¹ German lesson! The German word for contingency, "zusammenhängen", means "hang together".

Author: DataCamp / Andrew Bray & Lyon College / Marcus Birkenkrahe

Created: 2022-04-20 Wed 14:45